



Nested source-only formula selection (higher is better)




Each cell: 3 inner validation protocols x 2 seeds; 250 updates/model. Target development is excluded from its column.
Equal relation x destination-kind strata; log(n+2) - NLL. 32 raw-hash-ranked dev messages/protocol. Not test performance.

Final probability assigned to the correct destination



Selected per held-out target: DNS=hybrid, MODBUS=route_sink, TLS=hybrid, SMB2=hybrid
3-source -> 1-target; 600 updates, 3 fresh seeds. Historical test, exploratory. END/NULL are distinct destinations.
Modbus/TLS all END: high scores there do not by themselves establish arithmetic generalization.

Exact destination Hit@1





Selected per held-out target: DNS=hybrid, MODBUS=route_sink, TLS=hybrid, SMB2=hybrid
3-source -> 1-target; 600 updates, 3 fresh seeds. Historical test, exploratory. END/NULL are distinct destinations.
Uniform Hit@1 uses deterministic argmax (byte0), not random-sampling expected accuracy.

Per-field relation: final correct-destination probability



Selected per held-out target: DNS=hybrid, MODBUS=route_sink, TLS=hybrid, SMB2=hybrid
Mean over 3 seeds. DNS TXT has only 1 field; Modbus has 6. Decimal zeros are not substituted for missing support.

Paired final effects: selected formula versus current LAPA




Selected per held-out target: DNS=hybrid, MODBUS=route_sink, TLS=hybrid, SMB2=hybrid

95% Student-t intervals over 3 paired seed means (df=2), not packet-population confidence. Intervals are not clipped.


Formula choices were frozen before target inference; a negative result does not trigger reselection.

Locator, executor and final prediction: separate stage probabilities



Selected per held-out target: DNS=hybrid, MODBUS=route_sink, TLS=hybrid, SMB2=hybrid
Program probability is indexed at the true source only after a raw-only forward and prediction seal.
Source/Program are pre-validity heads; route probability additionally includes execution/validity normalization. No oracle inputs.

DNS ordinal-query shift is separate from the QKV hypothesis



DNS was held out. Source protocols have no positive training queries beyond slot3.
This is a conditional diagnostic subset, not a second training experiment. The fixed ordinal-query interface is unchanged.

Typed exact field-discovery F1 (separate from endpoint probability)

