


Final destination retrieval: Hit@1

Mean Hit@1, higher is better



3 source protocols -> 1 held-out target; three seeds; 600 updates. Shown: raw field means, not balanced selection scores.
P-selected / NLL-selected formulas vary by target (see selection table); no target-based reselection.
Any CNN/GRU champion is one canonical model reused across backbone panels. Tiny nonzero values use scientific notation.
Uniform Hit@1 uses deterministic argmax at byte 0, not sampled-hit expectation.